

LAB 2: SUPERVISED LEARNING – CLASSIFICATION

(Duration: 2 sessions)

(PART A: LOGISTIC REGRESSION)

Exercise 1: Implement Logistic Regression from Scratch

1. Implement the sigmoid function.
2. Write a class `LogisticRegression` with `fit`, `predict`, and `predict_proba` methods.
3. Train your model on a simple dataset (e.g., `make_classification()` or custom).
4. Compare your implementation with `sklearn.linear_model.LogisticRegression`.

Exercise 2 : Visualizing Logistic Regression Decision Boundaries

1. Train Logistic Regression on a 2D Dataset:
 - a) Use a simple 2D dataset such as `make_moons()` or `make_blobs()`.
 - b) Train a logistic regression model on this dataset.
2. Plot the Data and Decision Boundary:
 - a) Plot the training data.
 - b) Visualize the decision boundary of the logistic regression model.
3. Plot the Probability Map:
 - a) Use `predict_proba()` to plot the probability map of the classifier over the 2D space.
 - b) Color each data point by its predicted probability of belonging to class 1. Use a color gradient to highlight the model's confidence across the 2D space.
4. Effect of Regularization:
 - a) Train logistic regression with L2 regularization.
 - b) Visualize how the decision boundary changes with different values of regularization strength λ .
 - c) Discuss how regularization impacts the decision boundary.

Exercise 3 : Multiclass Logistic Regression with Softmax

1. Implement Multiclass Logistic Regression:
 - a) Using gradient descent, implement softmax regression from scratch.
 - b) Apply the model to the Iris dataset which has 3 classes.
2. Compare with `sklearn`'s Multinomial Logistic Regression:
 - a) Train a `sklearn.linear_model.LogisticRegression` with `multi_class='multinomial'` on the Iris dataset.
 - b) Compare the predictions from your softmax implementation and `sklearn`'s multinomial logistic regression.
 - c) Discuss any differences in implementation and performance.

Exercise 4 : Logistic Regression with Polynomial Features for Non-linear Data

1. Generate a Non-linear Dataset (Use `make_circles()` to generate a non-linearly separable dataset)
2. Apply Polynomial Features:
 - a) Use `PolynomialFeatures` to transform the data into a higher-dimensional space (quadratic terms).
 - b) Train a logistic regression model on the transformed features.
3. Visualize the Decision Boundary:
 - a) Plot the transformed data and the decision boundary of the logistic regression model after applying polynomial features.
 - b) Discuss how the transformation allows the model to handle non-linearly separable data.

(PART B: SUPPORT VECTOR MACHINES)

Exercise 5 : Hard-Margin SVM Implementation (Linearly Separable Case)

1. Generate a linearly separable dataset (e.g., `make_blobs`).
2. Formulate the primal problem:
$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{subject to} \quad y^{(i)}(\mathbf{w}^\top \mathbf{x}^{(i)} + b) \geq 1$$
3. Use an optimization package (`cvxopt`, `scipy.optimize`) to solve.
4. Visualize the decision boundary, support vectors, and margin.

Exercise 6: Soft-Margin SVM (Non-Separable Data)

1. Create the Dataset
 - a) Use `make_classification` to generate a 2D dataset with two overlapping classes. (Use `flip_y=0.1` and `class_sep=0.8` to add overlap)
 - b) Plot the data to check that it's not linearly separable.
2. Train a Soft-Margin SVM
 - a) Use `SVC(kernel='linear', C=1)` to train a linear SVM.
 - b) Plot the decision boundary and highlight the support vectors.
(What do you notice about the support vectors and the margin?)
3. Change the Regularization Parameter C
 - a) Train the SVM with different values of C: 0.01, 0.1, 1, 10, and 100.
 - b) Plot the decision boundaries for each case.
(Observe how the margin and number of misclassified points change.)
4. Evaluate the Model
 - a) For each value of C, calculate the accuracy on the training set.
 - b) Which value of C gives a good balance between accuracy and margin width?
 - c) What happens to the margin width when C increases?
 - d) If your data contains noise, would you prefer a large or small value for C? Why?
 - e) How can cross-validation help you choose the best value for C?

Exercise 7: Exploring Non-Linear SVM and Kernel Functions

1. Start with a Linearly Separable Dataset
 - a) Generate a simple linearly separable dataset using `make_classification()` or manual points.
 - b) Train a linear SVM and plot the decision boundary.
 - c) Calculate and display the training accuracy
 - d) Observe the margin and explain the classification results.
2. Try a Non-Linearly Separable Dataset
 - a) Generate a non-linear dataset using `make_moons()` or the XOR dataset.
 - b) Train a linear SVM and visualize the decision boundary.
 - c) Explain why the model fails to separate the classes.
3. Introduce and Apply the Polynomial Kernel
 - a) Train an SVM with a polynomial kernel on the previous non-linear dataset.
 - b) Experiment with different degrees d .
 - c) Plot and discuss how the decision boundary changes with increasing d .
 - d) Calculate and display the training accuracy
 - e) Compare results with the linear kernel.
4. Explore the RBF Kernel
 - a) Train an SVM using the RBF kernel on the same dataset.
 - b) Tune the parameter γ and observe its impact on the decision boundary.
 - c) Visualize the support vectors and the margin.
 - d) Calculate and display the training accuracy
5. Compare All Kernels
 - a) Train SVMs with linear, polynomial, and RBF kernels on the same dataset (e.g., `make_moons()` or XOR).
 - b) Plot all decision boundaries.
 - c) Discuss which kernel performed best and why.

Exercise 8: SVM on Real-World Data

1. Use the breast cancer dataset or Iris dataset.
2. Train with different kernels: linear, RBF, and polynomial.
3. Evaluate with accuracy, confusion matrix, and ROC (if binary).
4. Discuss when SVM performs best compared to other models.

Exercise 9: Multi-class SVM

1. Train SVC on the Iris dataset.
2. Understand One-vs-One (OvO) and One-vs-Rest (OvR) strategies.
3. Use `multi_class='ovr'` and `multi_class='ovo'` options.
4. Plot 2D decision regions after dimensionality reduction (e.g., PCA).